

D213_PA1

April 22, 2024

1 D212 PA 3 code - Doug Haunsperger

1.1 Data Preparation

1.1.1 Do initial package import and data read

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

df = pd.read_csv('medical_time_series.csv')
#view first 5 rows
df.head(5)
```

```
[1]:   Day  Revenue
0    1  0.000000
1    2 -0.292356
2    3 -0.327772
3    4 -0.339987
4    5 -0.124888
```

Set Day column to index

```
[2]: df.set_index('Day')
```

```
[2]:   Revenue
Day
1    0.000000
2   -0.292356
3   -0.327772
4   -0.339987
5   -0.124888
..      ...
727  15.722056
728  15.865822
729  15.708988
730  15.822867
731  16.069429
```


```
[731 rows x 1 columns]
```

1.1.2 C1. Line Graph

Here is a line graph of the realization of the input data set.

```
[3]: df.Revenue.plot()
```

```
[3]: <Axes: >
```



1.1.3 C2. Time Step Formatting

```
[4]: print(df.describe())  
      print(df.info())
```

	Day	Revenue
count	731.000000	731.000000
mean	366.000000	14.179608
std	211.165812	6.959905
min	1.000000	-4.423299
25%	183.500000	11.121742
50%	366.000000	15.951830

```

75%      548.500000    19.293506
max      731.000000    24.792249
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 731 entries, 0 to 730
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Day         731 non-null    int64
 1   Revenue     731 non-null    float64
dtypes: float64(1), int64(1)
memory usage: 11.6 KB
None

```

This time series is formatted as a daily chart, with 731 records and no gaps present in the data.

1.1.4 C3. Evaluate Stationarity

Stationarity - or lack of an overall trend in the data - is evaluated through the Augmented Dickey-Fuller test. The `adfuller()` function from the `statsmodels` package will accomplish this.

```

[5]: # Code ref: https://www.machinelearningplus.com/time-series/
      ↪ arima-model-time-series-forecasting-python/ (Prabhakaran, n.d.)
      from statsmodels.tsa.stattools import adfuller

      result = adfuller(df.Revenue)
      print('ADF Statistic: %f' % result[0])
      print('p-value: %f' % result[1])

```

```

ADF Statistic: -2.218319
p-value: 0.199664

```

P-value of 0.2 indicates that one cannot reject the null hypothesis that the series is non-stationary. Try a diff on the series and retest:

```

[6]: result = adfuller(df.Revenue.diff().dropna())
      print('ADF Statistic: %f' % result[0])
      print('p-value: %e' % result[1])

```

```

ADF Statistic: -17.374772
p-value: 5.113207e-30

```

Now there is a very significant result with a 1st order difference. Plot the resulting series:

```


[7]: df.Revenue.diff().dropna().plot()

```

```

[7]: <Axes: >

```



Create train-test split for out-of-time cross-validation and output train/test files.

Note - train/test is not randomized because “the order sequence of the time series should be intact in order to use it for forecasting.” (Prabhakaran, n.d.)

Note also that the output files are based on the non-differenced data since the ARIMA model will handle the differencing using a d parameter of $d=1$.

```
[8]: # Approx 2 years of data. Split into first 18 months (~ 547 days), last 6 months
train = df.Revenue[:547]
test = df.Revenue[547:]

# Also output differenced data
df_d = df.Revenue.diff().dropna().copy(deep=True)

train.to_csv("train.csv", index=False)
test.to_csv("test.csv", index=False)
df_d.to_csv("difference_1.csv", index=False)
```

1.2 Data Analysis & Model Identification

1.2.1 D1. Findings & Visualizations

Check for Seasonality, Trend; Plot decomposition

```
[9]: # Code ref: https://towardsdatascience.com/
      ↪finding-seasonal-trends-in-time-series-data-with-python-ce10c37aa861 (Hayes, ↪
      ↪2021).
      from statsmodels.tsa.seasonal import seasonal_decompose
      import datetime

      # Since the index is not a datetime object, I need to specify the periodicity ↪
      ↪of the data for seasonal_decompose. Loop through all values from weekly (7) ↪
      ↪to quarterly (91):
      max_seasonality_value = 0
      max_seasonality_period = 0
      for per in range(7,92,1):
          decomp_result = seasonal_decompose(df.Revenue, model="add", period=per)

          if np.max(decomp_result.seasonal) > max_seasonality_value:
              max_seasonality_value = np.max(decomp_result.seasonal)
              max_seasonality_period = per
              print(per, max_seasonality_value)

      decomp_result_max = seasonal_decompose(df.Revenue, model="add", ↪
      ↪period=max_seasonality_period)
      trend = decomp_result_max.trend
      seas = decomp_result_max.seasonal
      resids = decomp_result_max.resid


      decomp_result_max.plot();

      resids.plot()
```

```
7 0.03389461393396866
8 0.1026302002341985
14 0.25015445450146506
23 0.29420912411648636
25 0.3178594400012178
27 0.37732716662250726
28 0.37902884252646396
29 0.3843976535319445
30 0.45635891708796555
32 0.4777857659906336
37 0.5391212441615979
39 0.6245402903837007
43 0.6337591456761923
45 0.8205124164437018
49 0.9437095190724906
50 1.1159073435021423
53 1.278226410647288
```

```
78 1.3522448933152225
79 1.414494344519936
80 1.5120758610970633
```

```
[9]: <Axes: ylabel='Resid'>
```



The above analysis finds a slight “seasonal” signal at a period of 80 days, but the residuals (± 5) are much larger than the supposed seasonal adjustment of ± 1.5 . There is an upward trend in the revenue data for the first ~ 275 days, that then levels off for the remainder of the series.


By inspection, the residuals drift around the x-axis with no overall trend.


Autocorrelation Function

```
[10]: from statsmodels.tsa.stattools import acf, pacf
      from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

      plot_acf(df.Revenue, lags=90, title='Autocorrelation, original data')
      plt.show()

      plot_acf(df_d, lags=90, title='Autocorrelation, differenced data');
```






Spectral Analysis

```
[11]: # Code ref: https://ataspinar.com/2020/12/22/time-series-forecasting-with-stochastic-signal-analysis-techniques/
      ↪ (Taspinar, A., 2020)
y_vals = df.Revenue

fig, ax = plt.subplots(figsize=(14,3))
ax.plot(np.fft.rfft(y_vals))
ax.set_title('FFT of raw data')

plt.xlim(-1, 100)
plt.show()
```

```
C:\Users\dough\anaconda3\Lib\site-packages\matplotlib\cbook\__init__.py:1335:
ComplexWarning: Casting complex values to real discards the imaginary part
  return np.asarray(x, float)
```

There is a pronounced high value at low frequency; this is due to the overall trend in the raw data. I will now run the Fourier transform on the detrended data:

```
[12]: from scipy.signal import savgol_filter, find_peaks, periodogram

yvalues_trend = savgol_filter(y_vals,25,1)
yvalues_detrended = y_vals - yvalues_trend

fft_y_det = np.fft.fft(yvalues_detrended)
fft_y = np.abs(fft_y_det[:len(fft_y_det)//2])
indices_peaks = find_peaks(fft_y, height=80, distance=6)
indices_peaks = indices_peaks[0]
fft_x_ = np.fft.fftfreq(len(yvalues_detrended))
fft_x = fft_x_[:len(fft_x_)//2]

x_peaks = fft_x[indices_peaks]
y_peaks = fft_y[indices_peaks]

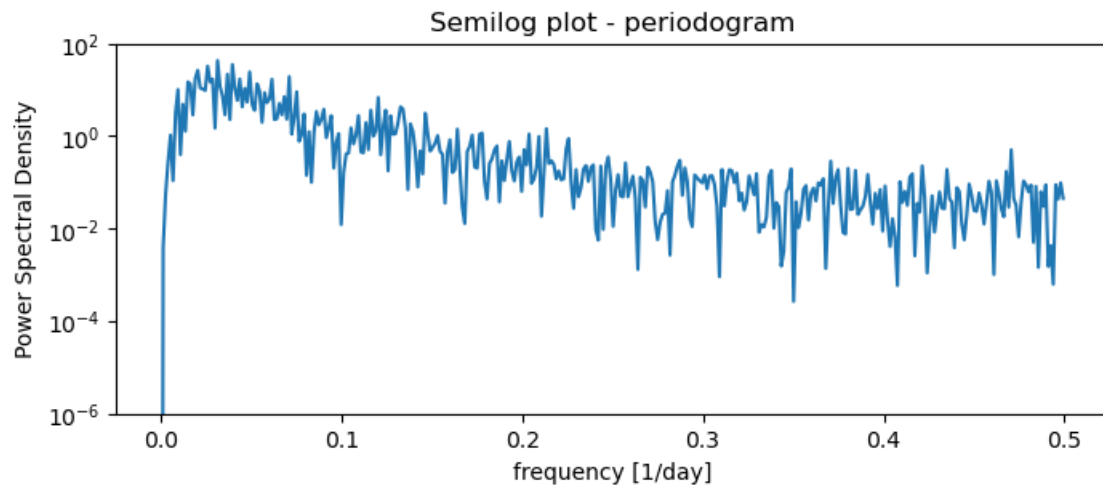
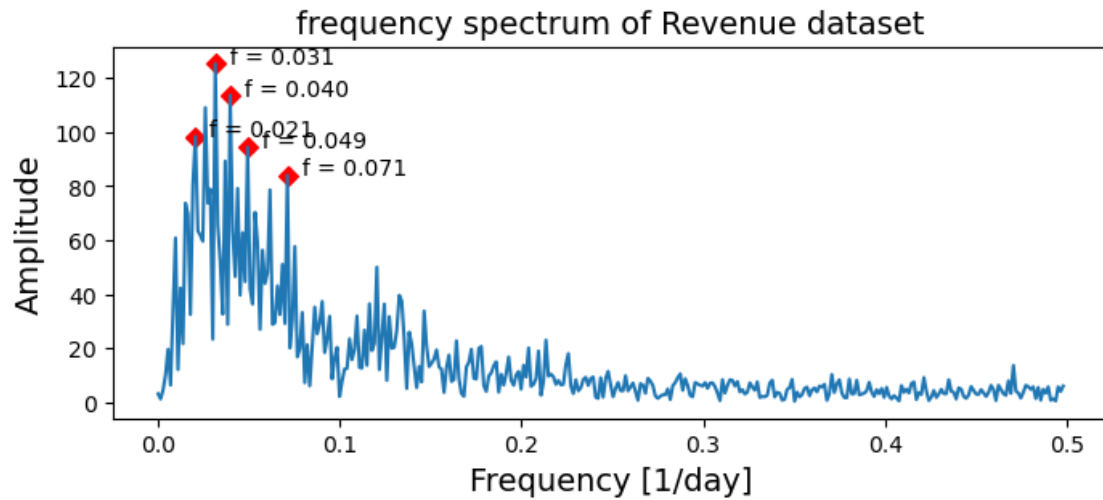
fig, ax = plt.subplots(figsize=(8,3))
ax.plot(fft_x, fft_y)
ax.scatter(x_peaks, y_peaks, color='red',marker='D')
ax.set_title('frequency spectrum of Revenue dataset', fontsize=14)
ax.set_ylabel('Amplitude', fontsize=14)
ax.set_xlabel('Frequency [1/day]', fontsize=14)

for idx in indices_peaks:
    x,y = fft_x[idx], fft_y[idx]
    text = " f = {:.3f}".format(x,y)
    ax.annotate(text, (x,y))

plt.show()

# Code ref: https://docs.scipy.org/doc/scipy/reference/generated/scipy.signal.
# periodogram.html#scipy.signal.periodogram (Scipy.org, 2024).
f, Pxx_den = periodogram(yvalues_detrended)
```

```
plt.figure(figsize=(8,3))
plt.semilogy(f, Pxx_den)
plt.ylim([1e-6,100])
plt.title('Semilog plot - periodogram')
plt.xlabel('frequency [1/day]')
plt.ylabel('Power Spectral Density');
```



I used the `find_peaks` function from `scipy` to look for peaks in the frequency spectrum of the de-trended revenue dataset. The plot shows a few small peaks, but when plotted on a semilog chart, these peaks disappear into the noise floor.


1.2.2 D2. ARIMA Model Fit

In the differencing analysis above, I already showed that the d parameter of the model should be 1, so it only remains to find the order of the AR term p and the MA term q .

Find p I use the Partial Autocorrelation plot:

```
[13]: fig, axes = plt.subplots(1, 2)
      axes[0].plot(df_d); axes[0].set_title('1st Differencing')
      axes[1].set(ylim=(0,5))
      plot_pacf(df_d, ax=axes[1]);
```

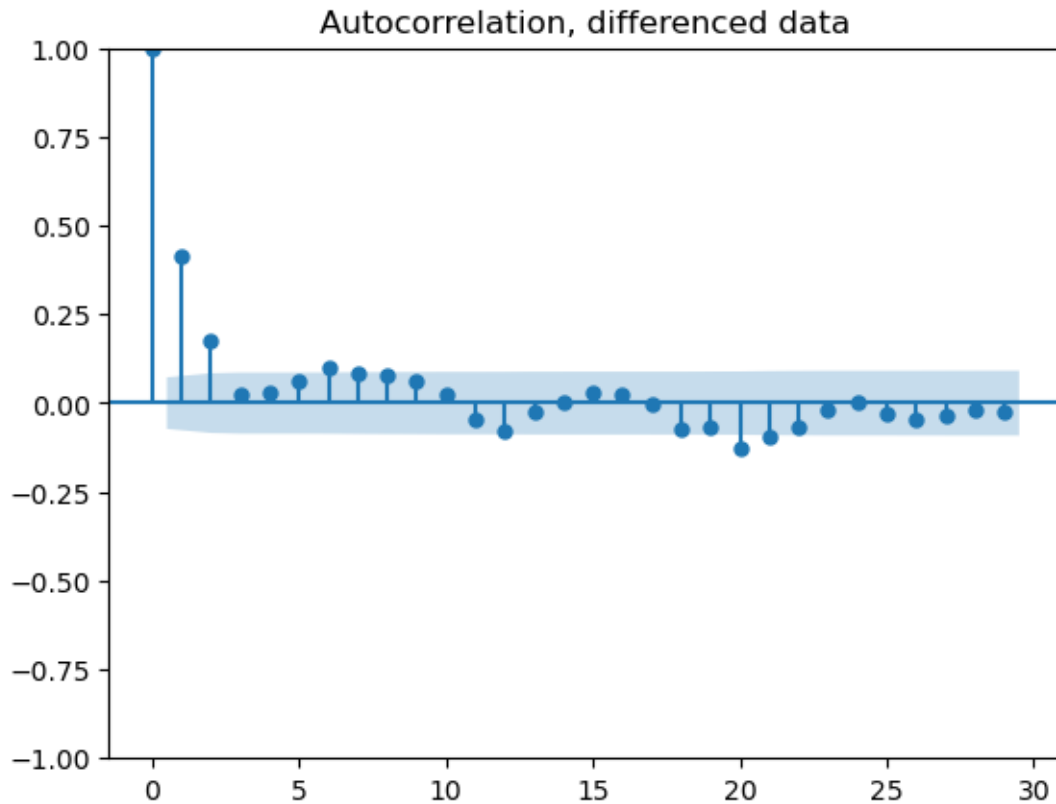
```
C:\Users\dough\anaconda3\Lib\site-packages\statsmodels\graphics\tsaplots.py:348:
FutureWarning: The default method 'yw' can produce PACF values outside of the
[-1,1] interval. After 0.13, the default will change to unadjusted Yule-Walker
('ywm'). You can use this method now by setting method='ywm'.
  warnings.warn(
```



Only the lag-1 partial autocorrelation is significantly different from 0, so that indicates the p AR order term should be 1.

Find q Recall the ACF of the differenced data looked like this:

```
[14]: plot_acf(df_d, title='Autocorrelation, differenced data');
```



The first 2 lags have a significant autocorrelation, which indicates that 2 is likely a good choice for q . This gives the (p,d,q) of the ARIMA model as $(1,1,2)$.

Fit an ARIMA model to the train data

```
[15]: from statsmodels.tsa.arima.model import ARIMA
      from statsmodels.graphics.tsaplots import plot_predict

      arima = ARIMA(train, trend='t', order=(1,1,2))
      result = arima.fit()

      result.summary()
```

```
[15]: <class 'statsmodels.iolib.summary.Summary'>
      """
                                     SARIMAX Results
=====
Dep. Variable:                Revenue    No. Observations:                547
Model:                        ARIMA(1, 1, 2)    Log Likelihood                -329.625
Date:                        Mon, 22 Apr 2024    AIC                        669.249
```

Time: 13:25:27 BIC 690.762
Sample: 0 HQIC 677.659
- 547

Covariance Type: opg

	coef	std err	z	P> z	[0.025	0.975]
x1	0.0220	0.032	0.690	0.490	-0.041	0.085
ar.L1	0.1200	0.191	0.627	0.531	-0.255	0.495
ma.L1	0.2771	0.189	1.464	0.143	-0.094	0.648
ma.L2	0.1839	0.073	2.506	0.012	0.040	0.328
sigma2	0.1958	0.013	15.147	0.000	0.170	0.221

===

Ljung-Box (L1) (Q): 0.00 Jarque-Bera (JB):
1.93

Prob(Q): 0.98 Prob(JB):
0.38

Heteroskedasticity (H): 1.08 Skew:
-0.03


Prob(H) (two-sided): 0.59 Kurtosis:
2.72

=====
===

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).
"""

```
[16]: fig, ax = plt.subplots()
df.Revenue.plot(ax=ax)
plot_predict(result, start=547, end=731, ax=ax)
test.plot(ax=ax, label='Test Revenue data')
plt.legend()
plt.show()
```



All the above is the procedure to find an ARIMA model manually. There exists an automated way using the pmdarima package. Let's check to see if it finds the same model or not.

```
[17]: #import sys
      #!{sys.executable} -m pip install pmdarima
      import pmdarima as pm
      automod = pm.auto_arima(train,start_p=0,start_q=0,test='adf', \
                             # Search for p & q between 0 and 4; use adfuller for d_
                             ↪determination
                             max_p=4,max_q=4, \
                             # No seasonality
                             seasonal=False, \
                             stepwise=False, trend=None,
                             trace=True, with_intercept=True)
```

```
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=767.498, Time=0.04 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=691.444, Time=0.04 sec
ARIMA(0,1,2)(0,0,0)[0] intercept : AIC=667.740, Time=0.08 sec
ARIMA(0,1,3)(0,0,0)[0] intercept : AIC=669.145, Time=0.07 sec
ARIMA(0,1,4)(0,0,0)[0] intercept : AIC=670.996, Time=0.10 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=670.883, Time=0.04 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=672.559, Time=0.07 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=669.249, Time=0.13 sec
```

```

ARIMA(1,1,3)(0,0,0)[0] intercept : AIC=670.767, Time=0.34 sec
ARIMA(1,1,4)(0,0,0)[0] intercept : AIC=672.676, Time=0.51 sec
ARIMA(2,1,0)(0,0,0)[0] intercept : AIC=672.407, Time=0.05 sec
ARIMA(2,1,1)(0,0,0)[0] intercept : AIC=671.328, Time=0.17 sec
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=670.349, Time=0.30 sec
ARIMA(2,1,3)(0,0,0)[0] intercept : AIC=672.708, Time=0.51 sec
ARIMA(3,1,0)(0,0,0)[0] intercept : AIC=669.366, Time=0.07 sec
ARIMA(3,1,1)(0,0,0)[0] intercept : AIC=671.221, Time=0.15 sec
ARIMA(3,1,2)(0,0,0)[0] intercept : AIC=671.969, Time=0.25 sec
ARIMA(4,1,0)(0,0,0)[0] intercept : AIC=671.097, Time=0.10 sec
ARIMA(4,1,1)(0,0,0)[0] intercept : AIC=670.145, Time=0.44 sec

```

Best model: ARIMA(0,1,2)(0,0,0)[0] intercept

Total fit time: 3.479 seconds

ARIMA(0,1,2) performs slightly better based on the AIC estimator value. I will rerun the model fit and prediction with no autoregressive component.

```

[18]: revised_arima = ARIMA(train, trend='t', order=(0,1,2))
      revised_result = revised_arima.fit()

      revised_result.summary()

```

```

[18]: <class 'statsmodels.iolib.summary.Summary'>

```

```

"""
                                SARIMAX Results
=====
Dep. Variable:                Revenue    No. Observations:                547
Model:                        ARIMA(0, 1, 2)    Log Likelihood                -329.870
Date:                        Mon, 22 Apr 2024    AIC                           667.740
Time:                        13:25:31          BIC                           684.951
Sample:                        0              HQIC                         674.468
                                - 547
Covariance Type:                opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
x1              0.0221      0.031        0.718      0.473      -0.038      0.082
ma.L1           0.3901      0.042        9.324      0.000       0.308      0.472
ma.L2           0.2173      0.042        5.189      0.000       0.135      0.299
sigma2          0.1959      0.013       15.296      0.000       0.171      0.221
=====
===
Ljung-Box (L1) (Q):                0.03    Jarque-Bera (JB):
1.80
Prob(Q):                0.86    Prob(JB):
0.41
Heteroskedasticity (H):            1.09    Skew:

```

```

-0.04
Prob(H) (two-sided):          0.58   Kurtosis:
2.73
=====
===

```

```

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-
step).
"""

```


```

[19]: fig, ax = plt.subplots()
df.Revenue.plot(ax=ax)
plot_predict(revised_result, start=547, end=731, ax=ax)
test.plot(ax=ax, label='Test Revenue data')
plt.legend()
plt.show()

forecast = revised_result.forecast(184)

mape = np.mean(np.abs(forecast - test)/np.abs(test))
print(mape)

```



0.15358598466166667

I get a mean absolute percentage error of 15.4% on the model prediction vs. the test data. Nearly all the test data is within the 95% confidence bounds.

1.2.3 D3. Forecast

Now I will refit the model on the full 731-record data set and forecast into the future.


```
[20]: forecast_mod = ARIMA(df.Revenue, trend='t', order=(0,1,2))
      fc = forecast_mod.fit()

      fig, ax = plt.subplots()
      df.Revenue.plot(ax=ax)
      plot_predict(fc, start=730, end=800, ax=ax)
      plt.xlim(700,800)
      plt.ylim(0,30)

      # Check for first prediction where CI is greater than $10M
      big_ci = 10
      ci=np.asarray(fc.get_prediction(start=730, end=790).conf_int(0.05))
      index_first=None
      for index, check in enumerate((ci[:,1]-ci[:,0]) > big_ci):
          if check == True:
              index_first = index
              break
      plt.axvline(x=730+index_first, color="g")
      print (index_first)

      plt.show()
```

14



After a forecast period of 14 days, the 95% confidence interval grows larger than \$10M, indicating increasing uncertainty in the forecast at longer time scales.